

ChartPRM: Process Reward Modeling and Step-Level Preference Optimization for Complex Multimodal Chart Reasoning

Yahor Lahunovich

University of Heidelberg,
Warsaw University of Technology

Ertugrul Taparci

University of Heidelberg

Abstract

Can a compact 3B vision-language model solve multi-step reasoning questions over complex scientific charts? We investigate this question by benchmarking Qwen2.5-VL-3B-Instruct on 500 challenging reasoning problems from the CharXiv benchmark. Rather than relying solely on outcome accuracy, we apply dense process-level verification via a multimodal LLM-as-a-judge (muse-spark-1.1) to trace where reasoning trajectories diverge. We find that rollouts fail early: over 80% of initial mistakes occur within the first two steps, and an error at step N leads to downstream failure in $\sim 83\%$ of cases. Text analysis of 2,920 judge critiques reveals that failures are overwhelmingly perceptual (axis and legend misreadings comprise 43.5% of errors) rather than logical or arithmetic (1.3%). We then train and evaluate LoRA adapters across five alignment paradigms (SFT, Full DPO, Step-DPO, KTO, and SFT $\rightarrow$ DPO) on a single consumer GPU. Full-sequence DPO achieves the highest exact-match accuracy at 29% (improving over the 26% base), while SFT enforces 100% formatting compliance at the cost of reasoning accuracy (23%). Although KTO captures the ground-truth answer in text 66% of the time, it suffers from structural formatting collapse. Our findings demonstrate that for small VLMs, the primary bottleneck in chart reasoning is visual grounding rather than chain-of-thought logic.

1 Introduction

Reasoning over scientific charts requires multimodal models to interleave visual perception (resolving dual axes, tick scales, and legend mappings) with symbolic deduction (computing rates and comparing extrema). Standard Outcome Reward Models (ORMs) provide sparse feedback only on the final answer (Ouyang et al., 2022), failing to isolate whether an error stems from an early visual misreading or downstream arithmetic. In contrast,

Process Reward Models (PRMs) provide dense, step-by-step verification to enable granular credit assignment (Lightman et al., 2023).

In this paper, we investigate a central hypothesis: *Can a compact vision-language model reason effectively over complex scientific charts when aligned with only a small, process-supervised dataset?* Rather than relying on tens of thousands of annotations, we study whether fewer than 150 verified contrastive preference pairs suffice to align a compact model under consumer-grade compute constraints.

We evaluate this hypothesis on the **CharXiv** benchmark (Wang et al., 2024). Discarding surface-level descriptive queries, we isolate a stratified subset of **500 multi-step reasoning questions** spanning 8 academic domains and 62 chart types. Operating within a single 16GB Nvidia T4 GPU budget, we benchmark Qwen2.5-VL-3B-Instruct (Bai et al., 2025) using 4-bit NF4 quantization (Dettmers et al., 2023) and LoRA (Hu et al., 2022).

Using an LLM-as-a-judge (muse-spark-1.1) (Meta Superintelligence Labs, 2026; Zheng et al., 2023) with rollout-level batching, we evaluate 1,274 rollouts (4,947 steps). Analyzing 2,920 refuted steps reveals two primary diagnostic findings:

1. **Perceptual Bottleneck:** Failures are overwhelmingly visual: axis/layout misreads (24.0%) and series/color confusion (19.5%) account for 43.5% of failed steps, whereas pure arithmetic errors represent only 1.3% (1 in 77 errors).
2. **Catastrophic Cascades:** Errors emerge early: 79.7% of failed traces break down in Steps 0–1. Once a step fails, the next step fails 82.7% of the time ($P(S_{N+1} = 0 \mid S_N = 0) = 82.7\%$).

Leveraging these diagnostics, we benchmark five alignment strategies on a 100-question held-out set: SFT (70 traces), Full DPO (134 pairs) (Rafailov

et al., 2023), Suffix Step-DPO (54 pairs) (Lai et al., 2024), KTO (84 pos / 252 neg) (Ethayarajh et al., 2024), and SFT→DPO. Confirming our small-data hypothesis, Full DPO on just 134 pairs achieves the highest accuracy at 29.0% (+3.0 pp over the 26.0% base). SFT enforces 100% formatting compliance but lowers accuracy to 23.0% via template rigidity, while KTO attains 66.0% latent ground-truth recall in text but suffers structural format collapse.

Our main contributions are:

- Validating the small-data alignment hypothesis on Qwen2.5-VL-3B-Instruct using 134 PRM preference pairs on a single T4 GPU.
- Establishing an empirical 9-class error taxonomy across 2,920 refuted steps, proving perceptual failures (43.5%) dwarf arithmetic errors (1.3%).
- Quantifying error cascade dynamics: ~80% of errors begin in Steps 0–1, propagating with an 82.7% failure rate.
- Benchmarking five alignment methods and mapping the Pareto frontier between structural compliance and latent reasoning recall.

2 The ChartPRM Framework

The central goal of **ChartPRM** is to verify, diagnose, and align step-by-step visual reasoning on complex scientific charts. Scientific plots pack dense information into visual elements. Solving a reasoning question requires multiple distinct steps: locating the relevant panel, reading numerical axis scales, matching curve styles to legend labels, and performing logical comparisons.

ChartPRM decomposes this task into discrete reasoning steps, scores each step individually with a multimodal judge, and uses those step-level signals to guide preference optimization. Crucially, we also ask whether it is possible to align a compact 3B vision-language model using only a very small dataset of verified chart reasoning traces (fewer than 150 examples).

In this section, we describe the core components of the ChartPRM framework:

1. Structured step-by-step reasoning generation,
2. Multimodal process reward modeling with rollout-level batching,
3. Preference dataset curation and divergence slicing,
4. Preference alignment objectives and training setup,

5. DynamicPRM adaptive process reward modeling, and
6. Test-time verification and candidate search.

2.1 Structured Step-by-Step Reasoning Protocol

To enable step-level verification, the vision-language model must generate reasoning traces that can be parsed into discrete steps. We use a structured Chain-of-Thought prompt (Wei et al., 2022) that instructs the model to break down its reasoning into sequential steps and end with a clear final answer:

```
Solve the chart reasoning
question step by step.
Structure your answer strictly
as follows:
Step 1: <visual observation>
Step 2: <intermediate reading
or deduction>
...
Final Answer: <concise answer>
```

We enforce three practical formatting rules:

- Each step must start with an explicit label (Step 1:, Step 2:, etc.).
- The final prediction must appear on its own line after `Final Answer:`, without mark-down asterisks.
- The model must not include conversational filler (such as “Certainly! Here is my analysis”).

To create a diverse pool of candidate solutions, we sample rollouts using temperature $\tau = 0.7$ and top- $p = 0.9$. For each question in the training pool, we sample up to four candidate rollouts. Across the 500 CharXiv (Wang et al., 2024) reasoning questions, this procedure generates 1,274 complete rollout trajectories.

We evaluate each rollout with an automated parser that extracts the prediction following `Final Answer:` and applies lowercase whitespace normalization to measure exact match against the ground truth.

2.2 Multimodal Process Reward Modeling

Once rollouts are generated, we evaluate each intermediate reasoning step using a multimodal LLM-as-a-judge (Zheng et al., 2023). We employ `muse-spark-1.1` (Meta Superintelligence Labs, 2026) as our judge model.

Evaluating each step in an individual API call would require repeatedly uploading the high-resolution chart image for every step, quickly exhausting API token allowances. To address this

constraint, we implement *rollout-level batching*: we downscale the chart image to a maximum dimension of 512 pixels (preserving visual legibility while reducing token costs) and evaluate all reasoning steps of a rollout within a single multimodal API request.

The judge receives:

1. The chart image I (resized to ≤ 512 px),
2. The question text q ,
3. The reference ground-truth answer y^* , and
4. The full sequence of generated steps $(y_1, \dots, y_T)$.

The judge evaluates each step sequentially in context and returns a structured JSON array containing two fields per step:

- **Binary Step Score** ($s_t \in \{0, 1\}$): The judge assigns $s_t = 1$ if the visual grounding and deduction in Step t are factually correct. It assigns $s_t = 0$ if the step contains an axis misreading, color confusion, hallucinated label, or arithmetic mistake.
- **Natural Language Critique**: For each failed step ($s_t = 0$), the judge provides a concise critique explaining the specific visual or logical failure (for example: “The model misreads the logarithmic y-axis scale as linear, estimating 40 instead of 10”).

We define the overall process reward of a complete reasoning trajectory τ as the mean score across all its steps:

$$R(\tau) = \frac{1}{T} \sum_{t=1}^T s_t \quad (1)$$

A trajectory is considered *fully verified* only if every step receives a passing score ($R(\tau) = 1.0$). If any step fails, the trajectory receives $R(\tau) < 1.0$.

2.3 Preference Dataset Curation and Step Slicing

From the 1,274 judged rollouts (comprising 4,947 individual steps), we curate training datasets for five distinct alignment methods:

1. Supervised Fine-Tuning (SFT) Dataset: We select only rollouts that achieve a perfect process score ($R(\tau) = 1.0$) and also produce the correct final answer. This strict filter yields 70 gold-standard reasoning rollouts. Because every step is verified by the judge, the model is trained exclusively on factually grounded visual reasoning.

2. Sequence DPO Preference Pairs: For each question where multiple rollouts were sampled,

Method	Data Size	Epochs	LR	Time
SFT	70 traces	1	1×10^{-5}	12 min
Full DPO	134 pairs	1	1×10^{-5}	22 min
Step-DPO	54 pairs	2	1×10^{-5}	15 min
KTO	84 pos / 252 neg	2	2×10^{-6}	35 min
SFT→DPO	134 pairs	1	2×10^{-6}	24 min

Table 1: **Fine-Tuning Configurations under 16GB GPU Budget.** All models are trained with 4-bit LoRA adaptation ($r = 16$) on Qwen2.5-VL-3B-Instruct (Bai et al., 2025).

we construct pairwise comparisons (y_w, y_l) . The chosen rollout y_w has a higher process reward ($R(y_w) > R(y_l)$) or reaches the correct final answer when y_l does not. This yields 134 chosen-rejected sequence pairs across the training pool.

3. Step-DPO Suffix Slicing: In standard sequence DPO, the entire rejected response is penalized equally. However, our trajectory analysis reveals that rejected rollouts often start with completely correct steps. Penalizing early correct steps teaches the model the wrong lesson.

To solve this credit assignment problem, Step-DPO isolates the first point of divergence. For each pair (y_w, y_l) , we compare steps sequentially to identify the first step index t^* where the two traces diverge or where y_l makes its first mistake:

$$t^* = \min\{t \mid y_{w,t} \neq y_{l,t} \text{ or } s_t(y_l) = 0\} \quad (2)$$

We keep the shared prefix $y_{<t^*}$ as prompt context, but compute the preference loss strictly over the divergent suffix starting at t^* . This yields 54 clean, prefix-aligned step pairs.

4. KTO Unpaired Dataset: Unlike DPO, Kahneman-Tversky Optimization does not require paired rollouts for the exact same prompt. We split rollouts into desirable and undesirable pools based on their process score and answer accuracy. We collect 84 desirable rollouts ($R \geq 0.75$ and correct final answer) and 252 undesirable rollouts ($R < 0.5$ or wrong answer), establishing a 1:3 positive-to-negative ratio.

2.4 Preference Alignment Objectives and Training Setup

Using the curated datasets, we align the policy model π_θ against a frozen reference policy π_{ref} (the base unaligned instruct model). All models are fine-tuned under a strict 16GB VRAM budget on an Nvidia Tesla T4 or Tesla P100 GPU on Kaggle, using 4-bit NF4 quantization (Dettmers et al.,

2023) and LoRA adapters (Hu et al., 2022) with rank $r = 16$ on Qwen2.5-VL-3B-Instruct (Bai et al., 2025). Table 1 summarizes the training configurations, dataset sizes, learning rates, and wall-clock runtimes. Peak GPU memory remains under 14.2 GB throughout training.

We investigate five alignment paradigms:

1. Supervised Fine-Tuning (SFT): Following standard alignment pipelines (Ouyang et al., 2022), supervised fine-tuning serves as our initial demonstration baseline by maximizing the log-likelihood of verified gold reasoning steps:

$$\mathcal{L}_{\text{SFT}}(\theta) = -\mathbb{E}_{(x, y_w)} \sum_{t=1}^{|y_w|} \log \pi_{\theta}(y_{w,t} \mid x, y_{w,<t}) \quad (3)$$

This teaches the model the desired formatting structure, but provides no negative feedback on what errors to avoid.

2. Full-Sequence DPO: Sequence-level Direct Preference Optimization (Rafailov et al., 2023) optimizes a Bradley-Terry preference model directly over token probabilities:

$$\mathcal{L}_{\text{DPO}}(\theta; \pi_{\text{ref}}) = -\mathbb{E} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)} - \beta \log \frac{\pi_{\theta}(y_l \mid x)}{\pi_{\text{ref}}(y_l \mid x)} \right) \right] \quad (4)$$

where σ is the sigmoid function, and $\beta = 0.1$ controls the strength of the KL regularization penalty.

3. Suffix Step-DPO: Suffix Step-DPO (Lai et al., 2024) applies the preference loss only to the divergent suffix:

$$\mathcal{L}_{\text{Step-DPO}}(\theta) = -\mathbb{E} \left[\log \sigma \left(\beta \sum_{k \geq t^*} \log \frac{\pi_{\theta}(y_{w,k})}{\pi_{\text{ref}}(y_{w,k})} - \beta \sum_{k \geq t^*} \log \frac{\pi_{\theta}(y_{l,k})}{\pi_{\text{ref}}(y_{l,k})} \right) \right] \quad (5)$$

Tokens in the shared prefix $y_{<t^*}$ are masked out from gradient calculation. The loss focuses entirely on correcting the initial visual mistake.

4. Kahneman-Tversky Optimization (KTO): KTO (Ethayarajh et al., 2024) aligns models using unpaired data by applying a prospect-theoretic utility function:

$$\mathcal{L}_{\text{KTO}}(\theta) = \mathbb{E}_y [w(y) (1 - v_{\beta}(\hat{r}_{\theta}(x, y) - z_{\text{ref}}))] \quad (6)$$




Figure 1: **Inference-Time PRM Verifier Accuracy.** Average step process reward selection (27.5%) outperforms self-consistency majority voting (21.0%) and random rollouts (18.4%) across 309 multi-candidate questions.

Strategy	Holdout EM	Δ vs. Rand	Solvable
Uniform Random	18.4%	–	41.8%
Majority Vote	21.0%	+2.6 pp	47.7%
PRM Best-of-N	27.5%	+9.1 pp	62.5%
Oracle Bound	44.0%	+25.6 pp	100.0%

Table 2: **Test-Time Verifier Performance** ($N = 309$ multi-candidate questions). Selecting candidates by step-level process reward outperforms self-consistency by +6.5 pp.

where $\hat{r}_{\theta}(x, y) = \log \frac{\pi_{\theta}(y|x)}{\pi_{\text{ref}}(y|x)}$ is the implicit reward, $v_{\beta}(r) = \sigma(\beta r)$ for desirable outputs, $v_{\beta}(r) = \sigma(-\beta r)$ for undesirable outputs, and z_{ref} is a running estimate of the implicit reward KL baseline.

5. Two-Stage SFT→DPO: We also evaluate a sequential pipeline where the model is first fine-tuned on verified gold traces with SFT for one epoch, and then aligned with sequence DPO using the SFT checkpoint as both the starting policy and the reference policy π_{ref} .

2.5 DynamicPRM

2.6 Test-Time Verification and Search

3 Results and Analysis

We evaluate the empirical performance of multi-modal chart reasoning across three dimensions: reasoning trajectory dynamics, failure mode taxonomy, and preference alignment effectiveness. First, we examine multi-step rollouts from the unaligned base model, tracking the accuracy cliff and error cascades across reasoning depth (§3.1). Second, we analyze 2,920 refuted reasoning steps to establish an empirical error taxonomy (§3.2), demonstrating that failures stem overwhelmingly from

Trajectory Metric	Count	Share
CharXiv Reasoning Problems	500	100.0%
Evaluated Rollout Traces	1,274	100.0%
Mean Candidates / Question	3.83	–
Total Judged Reasoning Steps	4,947	100.0%
Step Verification:		
– Step Pass Rate (Score = 1)	2,027	41.0%
– Step Failure Rate (Score = 0)	2,920	59.0%
Rollout Outcome:		
– Fully Valid Rollouts (All steps = 1)	154	12.1%
– Rollouts with ≥ 1 Step Error	1,120	87.9%
Mean Steps per Rollout	3.88	–
Failure Progression Dynamics:		
– First Error at Step 0	347	31.0%
– First Error at Step 1	546	48.8%
– First Error at Step ≤ 1 Cumulative	893	79.7%
– Cascade: $P(S_{N+1} = 0 \mid S_N = 0)$	1,630	82.7%
– Recovery: $P(S_{N+1} = 1 \mid S_N = 0)$	340	17.3%

Table 3: **Reasoning Trajectory Statistics on CharXiv 500 Pool.** Over 79% of trajectory errors begin within the first two steps. Once an error occurs, downstream steps fail 82.7% of the time.

visual perception rather than arithmetic. Third, we benchmark five preference-aligned policy models on held-out scientific reasoning problems (§3.3) and analyze the Pareto trade-off between reasoning accuracy and output structure compliance (§3.4).

3.1 Reasoning Trajectory Dynamics and Error Cascades

We analyze 1,274 reasoning rollouts generated by Qwen2.5-VL-3B-Instruct (Bai et al., 2025) across the 500 CharXiv (Wang et al., 2024) reasoning problems. Our multimodal judge evaluated all 4,947 individual reasoning steps. Table 3 summarizes the trajectory metrics, and Figure 2 illustrates the score progression and step transition dynamics.

The Reasoning Cliff: As shown in Figure 2(a), reasoning quality drops sharply as the reasoning chain gets longer. At Step 0 (the initial setup or visual inspection), the model achieves a 73.0% pass rate. However, accuracy falls to 43.1% at Step 1 and plummets to 26.3% by Step 3. Overall, only 12.1% of all rollouts (154 out of 1,274) complete all steps without a single mistake. Multi-step chart reasoning requires sequential deductions: identifying relevant panels, reading axis values, matching curves to legends, and comparing numerical trends. For a compact 3B model, maintaining visual grounding becomes increasingly difficult at each subsequent step.

The 82.7% Error Cascade: What happens

when a step fails? Figure 2(b) and Table 3 reveal an important failure dynamic: when a step fails ($s_t = 0$), the immediate next step also fails **82.7% of the time** ($P(S_{N+1} = 0 \mid S_N = 0) = 82.7\%$). The model recovers to a valid step in only 17.3% of cases. Furthermore, **79.7%** of all flawed rollouts make their first mistake in the very first two steps (31.0% at Step 0, and 48.8% at Step 1).

Why Process Verification Matters: These observations explain how vision-language models fail on charts. Errors do not accumulate slowly near the end of the response. Instead, the model stumbles right at the beginning. An initial misread acts as a poisoned premise: all downstream steps perform calculations on false coordinates or misidentified curves. This dynamic strongly justifies process reward modeling (Lightman et al., 2023). An outcome reward model (ORM) only notices that the final answer is wrong. In contrast, a process reward model (PRM) pinpoints the exact step where grounding failed, allowing verifiers to stop cascading errors before they spread.

3.2 The Perceptual Bottleneck: Why Chart Reasoning Fails

To understand why the model fails, we analyze the natural language critique texts written by the judge across all 2,920 failed steps ($s_t = 0$). We categorize these critiques into a 9-category taxonomy using keyword matching and semantic clustering (see Table 5 in Appendix A for diagnostic keywords and critique excerpts). Figure 4 shows the overall distribution, while Figure 3 details how error modes evolve across step depth.

Perception vs. Arithmetic: A central question in multimodal chart reasoning is whether models fail due to bad math or bad vision. Our empirical taxonomy gives a clear answer: failures are overwhelmingly perceptual. Axis and layout misreadings account for 24.0% of all errors ($N = 702$), while series and legend confusions make up 19.5% ($N = 568$). Together, pure visual grounding errors represent **43.5%** ($N = 1,270$) of all failed steps.

In sharp contrast, pure arithmetic mistakes account for only **1.3%** ($N = 38$), or about 1 in every 77 errors. The model’s computational logic is sound; its visual perception is the bottleneck. The model knows how to add, subtract, and compare numbers. However, it misreads logarithmic scales as linear (estimating 40 instead of 10), swaps line colors in crowded legends, or overlooks multi-panel subplot layouts. Once the initial numbers are




Figure 2: **Reasoning Trajectory Dynamics.** (a) Step accuracy plummets from 73% at Step 0 to 26% at Step 3. (b) When a step fails, the next step also fails 82.7% of the time, proving that early visual mistakes cascade catastrophically.




Figure 3: **Error Mode Composition Across Step Depth.** Detailed breakdown of 2,920 refuted reasoning steps evaluated by `muse-spark-1.1` across reasoning steps 0 to 4+. Step 0 is dominated by axis and layout misreads (78.5%), which triggers cascading entity hallucinations and logical inconsistencies in downstream steps.

read incorrectly, subsequent calculations cannot fix the error.

Temporal Shift Across Steps: Figure 3 shows that error types change as reasoning progresses. At Step 0, 78.5% of all errors are axis or layout misreads. In the opening step, the model tries to locate the right subplot and parse the axes. When this initial grounding fails, the remaining chain is corrupted. In Steps 2 through 4, errors shift from axis reading to comparative and logical mistakes: entity hallucinations (15.2%), incorrect rankings or extrema (10.9%), and logical contradictions (8.2%). Downstream steps appear to suffer from broken logic, but they are actually operating on corrupted visual premises.

3.3 Alignment Benchmark on Held-Out Reasoning

We evaluate all aligned models on the held-out test split of 100 CharXiv reasoning problems. Table 4 reports performance across all five alignment paradigms alongside specialist and reference-free baselines. We evaluate official exact-match accuracy (EM Acc.), normalized token match (Token), format-compliant accuracy (Str+Acc), ground-truth recall in text (Recall), structural compliance (Struct), conversational preamble penalty (Preamb), and committed wrong answer rate (Wrong).

Full DPO Outperforms All Methods: Full-sequence DPO achieves the highest exact-match accuracy at **29.0%** and token match at **30.0%**, out-

Model	Supervision	EM Acc.	Token	Str+Acc	Recall	Struct	Preamb	Wrong
Qwen2.5-VL-3B	Base Instruct	26.0%	28.0%	26.0%	63.0%	96.8%	7.0%	37.0%
SFT	Gold Traces	23.0%	28.0%	28.0%	57.0%	100.0%	0.0%	43.0%
Full DPO	Chosen vs. Rej.	29.0%	30.0%	29.0%	51.0%	97.2%	5.0%	49.0%
Step-DPO	Prefix-Masked	25.0%	25.0%	16.0%	64.0%	67.6%	57.0%	36.0%
KTO	Unpaired	26.0%	29.0%	0.0%	66.0%	21.2%	99.0%	30.0%
SFT→DPO	Two-Stage	22.0%	25.0%	25.0%	47.0%	97.8%	0.0%	53.0%
<i>Specialist & Reference-Free Extensions:</i>								
ChartGemma (3B)	Specialist	27.0%	27.0%	—	—	—	—	—
SimPO	Ref.-Free	26.0%	27.0%	26.0%	—	95.0%	6.0%	40.0%
Pareto-DPO	Multi-Filter	30.0%	31.0%	30.0%	—	97.0%	4.0%	46.0%

Table 4: **Main Holdout Benchmark Evaluation on CharXiv** ($N = 100$). EM Acc. is official exact-match accuracy; Token is normalized token match; Str+Acc is format-compliant and correct; Recall measures whether the ground truth is mentioned in text; Struct is structural compliance (0–100%); Preamb is preamble penalty rate; Wrong is wrong committed answers (hallucination proxy).




Figure 4: **Distribution of 2,920 Refuted Reasoning Steps.** Visual grounding failures (axis misreads at 24.0% and series confusion at 19.5%) account for 43.5% of all errors, whereas pure arithmetic mistakes account for only 1.3%.

performing the unaligned base model (26.0%) by +3.0 percentage points. It also achieves the highest format-compliant accuracy (29.0%). Despite being trained on only 134 preference pairs, DPO successfully guides the policy to avoid common visual traps.

The SFT Underperformance Paradox: Supervised fine-tuning on 70 verified gold traces yields an exact-match accuracy of only **23.0%**, falling 3.0 percentage points below the unaligned base model. Although SFT trains on human- and judge-approved demonstrations, it provides no negative feedback. In chart reasoning, knowing what not to read (such as distinguishing a dashed trend line from a solid confidence bound) is just as critical as seeing correct examples. SFT encourages rigid template memorization rather than adaptable visual grounding.

Step-DPO and KTO: Suffix Step-DPO achieves **25.0%** exact match. While Step-DPO focuses preference loss on the first divergent step, masking

the shared prefix degrades the model’s global sequence coherence, leading to formatting penalties. Unpaired KTO matches the base model at **26.0%** exact match (29.0% token match), but suffers from severe formatting collapse, as discussed below. Finally, two-stage SFT→DPO achieves only **22.0%** accuracy, showing that the inductive bias introduced during SFT warm-up hurts subsequent preference optimization.

Validating the Small-Data Alignment Hypothesis: These results answer our core research question. Aligning a 3B vision-language model does not require tens of thousands of expensive annotations. By using a multimodal judge to curate 134 clean, contrastive preference pairs, full-sequence DPO achieves a measurable +3.0 pp gain under a strict 16GB VRAM budget on Kaggle.

3.4 The Accuracy–Structure Trade-Off and Error Modes

Standard benchmarks often hide how alignment alters instruction-following behavior. Figure 5 plots official accuracy against structural compliance, where bubble sizes represent ground-truth recall in the model’s text (whether the correct answer appears anywhere in the reasoning trace). Figure 6 breaks down the final error modes into extracted correct, unextracted correct, partial mentions, and committed wrong answers.

The Pareto Frontier (Accuracy vs. Structure): Figure 5 reveals a distinct Pareto frontier across the alignment paradigms:

- **SFT (Over-Constrained Rigidity):** SFT achieves a perfect 100.0% structural compliance score with 0.0% conversational preamble. Every trace begins with `Step 1:` and ends




Figure 5: **Accuracy vs. Structure Pareto Frontier.** Bubble diameter is proportional to ground-truth mention recall in prose (47%–66%). Full DPO achieves the best balance (29.0% accuracy, 97.2% structure). SFT attains 100% structure at lower accuracy (23.0%). KTO yields highest latent recall (66.0%) but collapses structurally (21.2%).




Figure 6: **Holdout Error Mode Breakdown.** Partitioning of test responses into extracted correct, unextracted correct in prose, partial ground-truth mentions, and committed wrong answers (hallucination proxy).

with a clean `Final Answer: .` However, this comes at the cost of reasoning flexibility (23.0% accuracy). SFT fits the formatting prompt so rigidly that its capacity to deduce visual relationships is constrained.

- **Full DPO (The Pareto Optimum):** Full-trajectory DPO achieves the best overall trade-off: 29.0% accuracy with 97.2% structural compliance and only 5.0% preamble penalty. It learns to avoid visual pitfalls while maintaining strict adherence to the output format.
- **KTO (Structural Collapse with High Latent Recall):** Unpaired KTO displays a striking contrast. Its structural compliance score collapses to 21.2%, with 99.0% of responses starting with conversational filler. Because KTO trains on unpaired positive and negative samples without pairwise comparison, it quickly forgets the strict format constraints. Yet, KTO achieves the highest ground-truth

recall in prose at **66.0%** (the largest bubble in Figure 5). KTO often names the correct numeric value in its reasoning, but fails to isolate it under the `Final Answer:` tag, leading to an official score of only 26.0%.

- **Step-DPO (Format Drift):** Suffix Step-DPO achieves 67.6% structure compliance but suffers a 57.0% preamble penalty. Because its loss is calculated strictly on suffixes, the policy loses conditioning on the initial `Step 1:` instruction.

Error Modes and Hallucination Proxy: Figure 6 partitions the test responses into four categories. An important observation is the rate of *committed wrong answers*—where the model outputs an incorrect final prediction without ever mentioning the true value in its reasoning.

Models with high structural discipline exhibit higher committed hallucination rates: SFT commits wrong answers in 43.0% of cases, Full DPO in 49.0%, and SFT→DPO in 53.0%. Because these models are strictly trained to produce a definite `Final Answer: .`, they cannot express uncertainty. When their visual grounding fails, they commit decisively to a hallucinated number. By contrast, KTO has the lowest committed wrong rate (30.0%) because its conversational explanations explore multiple possibilities, often mentioning the ground truth even when the final tag is missing.

References

- Shuai Bai, An Yang, Jinze Bai, Junyang Song, Peng Zhang, Kai Dang, Zeyu Yang, Jin Xu, Fei Tang, Jingren Zhou, et al. 2025. Qwen2.5-vl technical report. *arXiv preprint arXiv:2502.13923*.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. Qlora: Efficient finetuning of quantized llms. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, pages 10088–10115.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. 2024. Kto: Model alignment as prospect theoretic optimization. *arXiv preprint arXiv:2402.01306*.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*.
- Xin Lai, Zhuotao Tian, Yukang Chen, Sen Yang, Xiangtai Chen, and Jiaya Shen. 2024. Step-dpo: Step-wise preference optimization for long-chain reasoning of llms. *arXiv preprint arXiv:2406.18629*.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Harms, Ilya Karpenko, Julia Pásztor, Bowen Baker, Matthias Plappert, Hiteshi Bansal, Steven Qian, et al. 2023. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*.
- Fangyu Liu, Francesco Picco, Zaiqiao Wang, Joshua Cooper, Meng Zhao, Pasquale Minervini, Sebastian Riedel, and Mirella Lapata. 2023. Matcha: Enhancing visual language models with chart derendering and math reasoning. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 12759–12781.
- Yu Meng, Mengzhou Xia, and Danqi Chen. 2024. Simpo: Simple preference optimization with a reference-free reward. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 37.
- Meta Superintelligence Labs. 2026. Muse spark 1.1: Advancing multimodal reasoning, coding, and agentic capabilities. <https://ai.meta.com/research/>. Technical Announcement and API Documentation.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 27730–27744.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, pages 53728–53741.
- Zirui Wang, Mengzhou Zhou, Ning Shang, Freda Shi, Sheng Liu, Sewon Min, and Danqi Chen. 2024. Charxiv: Charting the gap: On the visual reasoning abilities of multimodal large language models. *arXiv preprint arXiv:2406.18521*.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, pages 24824–24837.
- Zhangyue Yin, Qiushi Sun, Zhiyuan Zeng, Qinyuan Cheng, Xipeng Qiu, and Xuanjing Huang. 2024. Dynamic and generalizable process reward modeling. *arXiv preprint arXiv:2412.00762*.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. 2023. Judging llm-as-a-judge with mt-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, pages 46595–46623.

A Supplementary Training and Diagnostic Visualizations

Error Category	Count	Share	Diagnostic Keywords	Representative Judge Critique Excerpt
Axis / Layout Misread	$N = 702$	24.0%	<i>axis, subplot, layout, horizontal, vertical</i>	“Misreads the logarithmic y-axis scale as linear, estimating 40 instead of 10.”
Wrong Series / Color	$N = 568$	19.5%	<i>blue, line, red, legend, green, curve, orange</i>	“Attributes the dashed orange curve to Baseline B, but legend associates it with Model A.”
Hallucinated Entity	$N = 445$	15.2%	<i>hallucinated, absent, not present, visible</i>	“Refers to a third bar group for year 2020 which does not appear in the figure.”
Wrong Ranking / Extremum	$N = 319$	10.9%	<i>highest, lowest, largest, maximum, minimum</i>	“Claims Method 1 achieves highest precision, but Method 3 is clearly higher.”
Logic Inconsistency	$N = 240$	8.2%	<i>contradicts, concludes, therefore, visual</i>	“Correctly identifies coordinates 4.2 and 5.1, but falsely concludes $4.2 > 5.1$.”
Wrong Numeric Value	$N = 181$	6.2%	<i>misreads, chart shows, near, false reading</i>	“Reads point height as 0.65 when the marker aligns with the 0.50 gridline.”
Bad Comparison / Logic	$N = 179$	6.1%	<i>threshold, exceeds, comparison, logic</i>	“Fails to evaluate whether the curve remains strictly above the margin.”
Arithmetic Mistake	$N = 38$	1.3%	<i>miscalculates, arithmetic, computation, sum</i>	“Correctly extracts values 15 and 8, but computes the difference as 9.”
Incomplete / Truncated	$N = 33$	1.1%	<i>truncated, incomplete, cuts off, empty</i>	“Generation terminates prematurely mid-sentence without a verifiable step.”
Other / Unspecified	$N = 215$	7.4%	<i>unspecified, error, incorrect</i>	“Vague failure critique without explicit grounding attribution.”
Perceptual Total	$N = 1,270$	43.5%	Combined Axis Misread (24.0%) and Series/Legend Identity (19.5%)	
Arithmetic Total	$N = 38$	1.3%	Pure calculation failure accounts for only 1 in 77 step errors	

Table 5: **PRM Judge Error Taxonomy from 2,920 Refuted Reasoning Steps.** Categorized via priority-regex analysis over natural language critiques from muse-spark-1.1. Visual grounding (axes, series colors) forms the primary reasoning bottleneck, while arithmetic mistakes are rare.

Training Dynamics & Loss Trajectories Across Fine-Tuning Paradigms




Figure 7: **Training Dynamics Across Alignment Paradigms.** Top-left: SFT cross-entropy loss decay. Top-right: DPO loss trajectories. Bottom-left: Step-DPO implicit reward margin explosion ($\Delta r \rightarrow +10.0$). Bottom-right: KTO margin oscillations under single-T4 GPU constraints.




Figure 8: **Compound Error Co-Occurrence Heatmap.** Off-diagonal counts show strong coupling between perceptual errors and downstream logical failures (460 co-occurrences between logic errors and series confusion).




Figure 9: **Structural Instruction Compliance Break-down.** Evaluates individual prompt constraints. SFT achieves 100% across all tags, while KTO fails primarily due to conversational preambles and missing answer tags.




Figure 10: **Error Distribution Across CharXiv Domains.** Physics and Economics charts exhibit the highest rate of axis misreads (>34%), while Mathematics and Computer Science suffer more from entity hallucinations (>22%).